



Experiment No. 6
Apply Boosting Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:20-09-2023
Date of Submission:09-10-2023



Aim: Apply Boosting algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Apply Boosting algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

Theory:

Suppose that as a patient, you have certain symptoms. Instead of consulting one doctor, you choose to consult several. Suppose you assign weights to the value or worth of each doctor's diagnosis, based on the accuracies of previous diagnosis they have made. The final diagnosis is then a combination of the weighted diagnosis. This is the essence behind boosting.

Algorithm: Adaboost- A boosting algorithm—create an ensemble of classifiers. Each one gives a weighted vote.

Input:

- D , a set of d class labelled training tuples
- k , the number of rounds (one classifier is generated per round)
- a classification learning scheme

Output: A composite model

Method

1. Initialize the weight of each tuple in D is $1/d$
2. For $i=1$ to k do // for each round
3. Sample D with replacement according to the tuple weights to obtain D_i
4. Use training set D_i to derive a model M_i
5. Compute $\text{error}(M_i)$, the error rate of M_i
6. $\text{Error}(M_i) = \sum w_j \cdot \text{err}(X_j)$
7. If $\text{Error}(M_i) > 0.5$ then
8. Go back to step 3 and try again
9. endif
10. for each tuple in D_i that was correctly classified do
11. Multiply the weight of the tuple by $\text{error}(M_i)/(1-\text{error}(M_i))$
12. Normalize the weight of each tuple
13. end for

To use the ensemble to classify tuple X

1. Initialize the weight of each class to 0



2. for $i=1$ to k do // for each classifier
3. $w_i = \log((1 - \text{error}(M_i)) / \text{error}(M_i))$ // weight of the classifiers vote
4. $C = M_i(X)$ // get class prediction for X from M_i
5. Add w_i to weight for class C
6. end for
7. Return the class with the largest weight.

Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

exp-6-ml

October 9, 2023

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import io
from sklearn.metrics import accuracy_score, precision_score, f1_score,
    ↪confusion_matrix, classification_report
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error
```

```
[ ]: df = pd.read_csv('./adult.csv')
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    32561 non-null  int64
1   workclass              32561 non-null  object
2   fnlwgt                 32561 non-null  int64
3   education              32561 non-null  object
4   education.num          32561 non-null  int64
5   marital.status         32561 non-null  object
6   occupation             32561 non-null  object
7   relationship           32561 non-null  object
8   race                   32561 non-null  object
9   sex                    32561 non-null  object
10  capital.gain            32561 non-null  int64
11  capital.loss            32561 non-null  int64
12  hours.per.week          32561 non-null  int64
13  native.country          32561 non-null  object
14  income                  32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
[ ]: df.head(5)
```

```
[ ]:
  age workclass  fnlwgt   education  education.num marital.status \
0   90      ?    77053      HS-grad             9      Widowed
1   82  Private  132870      HS-grad             9      Widowed
2   66      ?   186061  Some-college            10      Widowed
3   54  Private  140359      7th-8th             4      Divorced
4   41  Private  264663  Some-college            10      Separated

      occupation  relationship   race   sex  capital.gain \
0              ?  Not-in-family  White  Female           0
1  Exec-managerial  Not-in-family  White  Female           0
2              ?    Unmarried  Black  Female           0
3  Machine-op-inspct    Unmarried  White  Female           0
4   Prof-specialty    Own-child  White  Female           0

  capital.loss  hours.per.week  native.country  income
0          4356             40  United-States  <=50K
1          4356             18  United-States  <=50K
2          4356             40  United-States  <=50K
3          3900             40  United-States  <=50K
4          3900             40  United-States  <=50K
```

```
[ ]: df.describe()
```

```
[ ]:
count    32561.000000  3.256100e+04  32561.000000  32561.000000  32561.000000 \
mean      38.581647  1.897784e+05     10.080679   1077.648844     87.303830
std       13.640433  1.055500e+05     2.572720   7385.292085    402.960219
min       17.000000  1.228500e+04     1.000000     0.000000     0.000000
25%       28.000000  1.178270e+05     9.000000     0.000000     0.000000
50%       37.000000  1.783560e+05    10.000000     0.000000     0.000000
75%       48.000000  2.370510e+05    12.000000     0.000000     0.000000
max       90.000000  1.484705e+06    16.000000   99999.000000   4356.000000

      hours.per.week
count    32561.000000
mean      40.437456
std       12.347429
min        1.000000
25%       40.000000
50%       40.000000
75%       45.000000
max       99.000000
```

```
[ ]: for i in df.columns:
      t = df[i].value_counts()
```

```

index = list(t.index)
print ("Count of ? in", i)
for i in index:
    temp = 0
    if i == '?':
        print (t['?'])
        temp = 1
        break
if temp == 0:
    print ("0")

```

```

Count of ? in age
0
Count of ? in workclass
1836
Count of ? in fnlwgt
0
Count of ? in education
0
Count of ? in education.num
0
Count of ? in marital.status
0
Count of ? in occupation
1843
Count of ? in relationship
0
Count of ? in race
0
Count of ? in sex
0
Count of ? in capital.gain
0
Count of ? in capital.loss
0
Count of ? in hours.per.week
0
Count of ? in native.country
583
Count of ? in income
0

```

```

[ ]: df=df.loc[(df['workclass'] != '?') & (df['native.country'] != '?')]
print(df.head())

```

	age	workclass	fnlwgt	education	education.num	marital.status	\
1	82	Private	132870	HS-grad	9	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	

4	41	Private	264663	Some-college	10	Separated
5	34	Private	216864	HS-grad	9	Divorced
6	38	Private	150601	10th	6	Separated

	occupation	relationship	race	sex	capital.gain	\
1	Exec-managerial	Not-in-family	White	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	
5	Other-service	Unmarried	White	Female	0	
6	Adm-clerical	Unmarried	White	Male	0	

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	<=50K
3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K
5	3770	45	United-States	<=50K
6	3770	40	United-States	<=50K

```
[ ]: df["income"] = [1 if i=='>50K' else 0 for i in df["income"]]
print(df.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
1	82	Private	132870	HS-grad	9	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	
5	34	Private	216864	HS-grad	9	Divorced	
6	38	Private	150601	10th	6	Separated	

	occupation	relationship	race	sex	capital.gain	\
1	Exec-managerial	Not-in-family	White	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	
5	Other-service	Unmarried	White	Female	0	
6	Adm-clerical	Unmarried	White	Male	0	

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	0
3	3900	40	United-States	0
4	3900	40	United-States	0
5	3770	45	United-States	0
6	3770	40	United-States	0

```
[ ]: df_more=df.loc[df['income'] == 1]
print(df_more.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
7	74	State-gov	88638	Doctorate	16	Never-married	
10	45	Private	172274	Doctorate	16	Divorced	

11	38	Self-emp-not-inc	164526	Prof-school	15	Never-married
12	52	Private	129177	Bachelors	13	Widowed
13	32	Private	136204	Masters	14	Separated


	occupation	relationship	race	sex	capital.gain \
7	Prof-specialty	Other-relative	White	Female	0
10	Prof-specialty	Unmarried	Black	Female	0
11	Prof-specialty	Not-in-family	White	Male	0
12	Other-service	Not-in-family	White	Female	0
13	Exec-managerial	Not-in-family	White	Male	0

	capital.loss	hours.per.week	native.country	income
7	3683	20	United-States	1
10	3004	35	United-States	1
11	2824	45	United-States	1
12	2824	20	United-States	1
13	2824	55	United-States	1


```
[ ]: workclass_types = df_more['workclass'].value_counts()
labels = list(workclass_types.index)
aggregate = list(workclass_types)
print(workclass_types)
print(aggregate)
print(labels)
```

```
Private          4876
Self-emp-not-inc  714
Local-gov        609
Self-emp-inc     600
Federal-gov      365
State-gov        344
Name: workclass, dtype: int64
[4876, 714, 609, 600, 365, 344]
['Private', 'Self-emp-not-inc', 'Local-gov', 'Self-emp-inc', 'Federal-gov',
'State-gov']
```

```
[ ]: plt.figure(figsize=(7,7))
plt.pie(aggregate, labels=labels, autopct='%1.1f%%', shadow = True)
plt.axis('equal')
plt.show()
```

```
[ ]: #Count plot on single categorical variable
sns.countplot(x='income', data = df)
plt.show()
df['income'].value_counts()
```




```
[ ]: 0    22661  
     1     7508  
     Name: income, dtype: int64
```

```
[ ]: #Plot figsize  
     plt.figure(figsize=(10,7))  
     sns.heatmap(df.corr(), cmap='coolwarm', annot=True)  
     print(plt.show())
```

<ipython-input-16-6201d8194dba>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
     sns.heatmap(df.corr(), cmap='coolwarm', annot=True)
```



None

```
[ ]: plt.figure(figsize=(10,7))
sns.distplot(df['age'], color="red", bins=100)
plt.ylabel("Distribution", fontsize = 10)
plt.xlabel("Age", fontsize = 10)
plt.show()
```


<ipython-input-17-1b72b8b67fa9>:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['age'], color="red", bins=100)
```



```
[ ]: #To find distribution of categorical columns w.r.t income
fig, axes = plt.subplots(figsize=(20, 10))
plt.subplot(231)
sns.countplot(x='workclass',
              hue='income',
              data = df,
              palette="BuPu")
plt.xticks(rotation=90)
plt.subplot(232)
sns.countplot(x='marital.status',
              hue='income',
              data = df,
              palette="deep")
plt.xticks(rotation=90)
plt.subplot(233)
sns.countplot(x='education',
              hue='income',
              data = df,
              palette = "autumn")
plt.xticks(rotation=90)
plt.subplot(234)
sns.countplot(x='relationship',
```

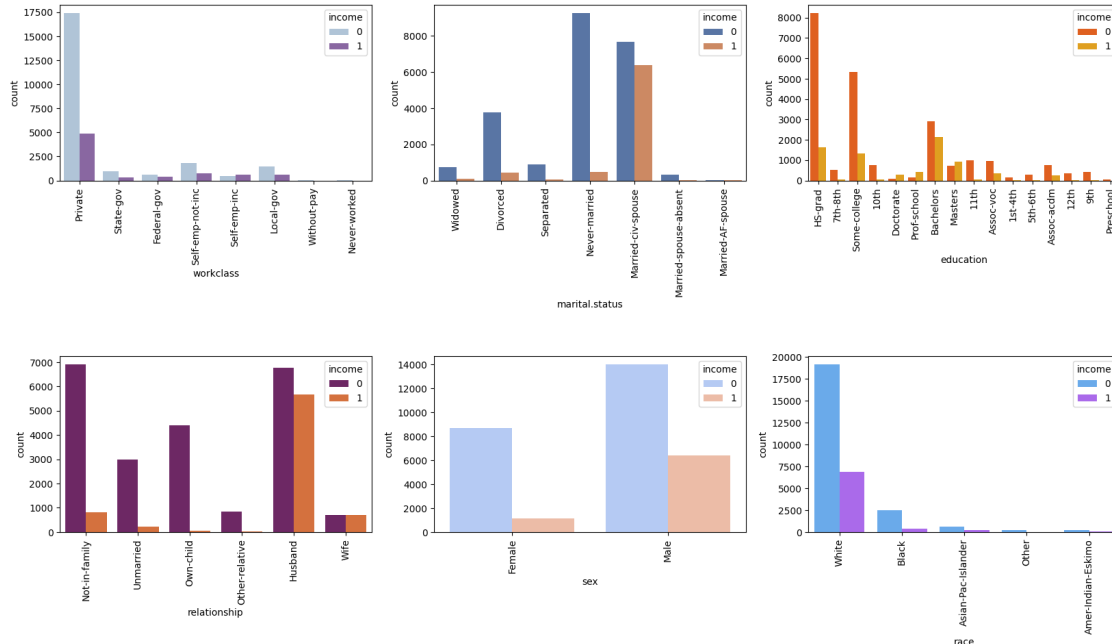
```

        hue='income',
        data = df,
        palette = "inferno")
plt.xticks(rotation=90)
plt.subplot(235)
sns.countplot(x='sex',
              hue='income',
              data = df,
              palette = "coolwarm")
plt.xticks(rotation=90)
plt.subplot(236)
sns.countplot(x='race',
              hue='income',
              data = df,
              palette = "cool")
plt.xticks(rotation=90)
plt.subplots_adjust(hspace=1)
plt.show()

```

<ipython-input-19-c1b6c6ef45b7>:3: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.

```
plt.subplot(231)
```



```
[ ]: df1 = df.copy()
```

```
[ ]: categorical_features = list(df1.select_dtypes(include=['object']).columns)
print(categorical_features)
df1
```

```
['workclass', 'education', 'marital.status', 'occupation', 'relationship',
'race', 'sex', 'native.country']
```

```
[ ]:
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
1	82	Private	132870	HS-grad	9	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	
5	34	Private	216864	HS-grad	9	Divorced	
6	38	Private	150601	10th	6	Separated	
...	
32556	22	Private	310152	Some-college	10	Never-married	
32557	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	
32558	40	Private	154374	HS-grad	9	Married-civ-spouse	
32559	58	Private	151910	HS-grad	9	Widowed	
32560	22	Private	201490	HS-grad	9	Never-married	

	occupation	relationship	race	sex	capital.gain	\
1	Exec-managerial	Not-in-family	White	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	
5	Other-service	Unmarried	White	Female	0	
6	Adm-clerical	Unmarried	White	Male	0	
...	
32556	Protective-serv	Not-in-family	White	Male	0	
32557	Tech-support	Wife	White	Female	0	
32558	Machine-op-inspct	Husband	White	Male	0	
32559	Adm-clerical	Unmarried	White	Female	0	
32560	Adm-clerical	Own-child	White	Male	0	

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	0
3	3900	40	United-States	0
4	3900	40	United-States	0
5	3770	45	United-States	0
6	3770	40	United-States	0
...
32556	0	40	United-States	0
32557	0	38	United-States	0
32558	0	40	United-States	1
32559	0	40	United-States	0
32560	0	20	United-States	0

```
[30169 rows x 15 columns]
```

```
[ ]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
for feat in categorical_features:
    df1[feat] = le.fit_transform(df1[feat].astype(str))
df1
```

```
[ ]:      age  workclass  fnlwgt  education  education.num  marital.status  \
1      82          3  132870         11           9           6
3      54          3  140359          5           4           0
4      41          3  264663         15          10           5
5      34          3  216864         11           9           0
6      38          3  150601          0           6           5
...  ...      ...      ...      ...      ...      ...
32556  22          3  310152         15          10           4
32557  27          3  257302          7          12           2
32558  40          3  154374         11           9           2
32559  58          3  151910         11           9           6
32560  22          3  201490         11           9           4

      occupation  relationship  race  sex  capital.gain  capital.loss  \
1              4              1    4    0           0          4356
3              7              4    4    0           0          3900
4             10              3    4    0           0          3900
5              8              4    4    0           0          3770
6              1              4    4    1           0          3770
...      ...      ...      ...      ...      ...
32556          11              1    4    1           0           0
32557          13              5    4    0           0           0
32558           7              0    4    1           0           0
32559           1              4    4    0           0           0
32560           1              3    4    1           0           0

      hours.per.week  native.country  income
1              18              38      0
3              40              38      0
4              40              38      0
5              45              38      0
6              40              38      0
...      ...      ...      ...
32556          40              38      0
32557          38              38      0
32558          40              38      1
32559          40              38      0
32560          20              38      0
```

[30169 rows x 15 columns]

```
[ ]: X = df1.drop(columns = ['income'])
y = df1['income'].values
# Splitting the data set into train and test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.
↳3, random_state = 0)
print ("Train set size: ", X_train.shape)
print ("Test set size: ", X_test.shape)
```

Train set size: (21118, 14)

Test set size: (9051, 14)

```
[ ]: from sklearn.ensemble import AdaBoostClassifier
# Train Adaboost Classifier
abc = AdaBoostClassifier(n_estimators = 300, learning_rate=1)
abc_model = abc.fit(X_train, y_train)
#Prediction
y_pred_abc = abc_model.predict(X_test)
```


```
[ ]: print("Accuracy: ", accuracy_score(y_test, y_pred_abc))
print("F1 score :", f1_score(y_test, y_pred_abc, average='binary'))
print("Precision : ", precision_score(y_test, y_pred_abc))
```

Accuracy: 0.8637719588995691

F1 score : 0.7008007765105557

Precision : 0.7921009325287987

```
[ ]: cm = confusion_matrix(y_test, y_pred_abc)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap=
↳"coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:' +
↳str(round(accuracy_score(y_test,y_pred_abc), 2)), size = 15);
plt.show()
print(classification_report(y_test, y_pred_abc))
```

	precision	recall	f1-score	support
0	0.88	0.94	0.91	6753
1	0.79	0.63	0.70	2298
accuracy			0.86	9051
macro avg	0.84	0.79	0.81	9051
weighted avg	0.86	0.86	0.86	9051

```
[ ]: from sklearn.ensemble import GradientBoostingClassifier
      #Training the model with gradient boosting
      gbc = GradientBoostingClassifier(
      learning_rate = 0.1,
      n_estimators = 500,
      max_depth = 5,
      subsample = 0.9,
      min_samples_split = 100,
      max_features='sqrt',
      random_state=10)
      gbc.fit(X_train,y_train)
      # Predictions
```

```

y_pred_gbc = gbc.predict(X_test)
print("Accuracy : ",accuracy_score(y_test, y_pred_gbc))
print("F1 score : ", f1_score(y_test, y_pred_gbc, average = 'binary'))
print("Precision : ", precision_score(y_test, y_pred_gbc))

```

```

Accuracy :  0.8689647552756602
F1 score :  0.7218574108818011
Precision : 0.7828077314343845

```

```

[ ]: rms = np.sqrt(mean_squared_error(y_test, y_pred_gbc))
print("RMSE for gradient boost: ", rms)

```

```


RMSE for gradient boost:  0.3619879068758235

```

```

[ ]: cm = confusion_matrix(y_test, y_pred_gbc)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot = True, fmt=".3f", linewidths = 0.5, square = True, cmap=
    ↪"coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:' + str(round(accuracy_score(y_test,
    ↪y_pred_gbc),2)), size = 15);
plt.show()
print(classification_report(y_test, y_pred_gbc))

```



	precision	recall	f1-score	support
0	0.89	0.94	0.91	6753
1	0.78	0.67	0.72	2298
accuracy			0.87	9051
macro avg	0.84	0.80	0.82	9051
weighted avg	0.86	0.87	0.87	9051

```
[ ]: import xgboost as xgb
from xgboost import XGBClassifier
#Training the model with gradient boosting
xgboost = XGBClassifier(learning_rate=0.01,
colsample_bytree = 0.4,
n_estimators=1000,
max_depth=20,
gamma=1)
xgboost_model = xgboost.fit(X_train, y_train)
```


```
[ ]: # Predictions
y_pred_xgboost = xgboost_model.predict(X_test)
print("Accuracy : ",accuracy_score(y_test, y_pred_xgboost))
print("F1 score : ", f1_score(y_test, y_pred_xgboost, average = 'binary'))
print("Precision : ", precision_score(y_test, y_pred_xgboost))
```

```
Accuracy : 0.8655397193680257
F1 score : 0.7090604829070045
Precision : 0.786737400530504
```

```
[ ]: rms = np.sqrt(mean_squared_error(y_test, y_pred_xgboost))
print("RMSE for xgboost: ", rms)
```


```
RMSE for xgboost: 0.3666882608319693
```

```
[ ]: cm = confusion_matrix(y_test, y_pred_xgboost)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap_
↪="coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:
↪'+str(round(accuracy_score(y_test,y_pred_xgboost),2)), size = 15);
plt.show()
print(classification_report(y_test,y_pred_xgboost))
```



	precision	recall	f1-score	support
0	0.89	0.94	0.91	6753
1	0.79	0.65	0.71	2298
accuracy			0.87	9051
macro avg	0.84	0.79	0.81	9051
weighted avg	0.86	0.87	0.86	9051

```
[ ]: from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_test, xgboost.predict_proba(X_test)[: ,1])
plt.figure(figsize = (10,5))
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr, tpr)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC CURVE for Xgboost')
plt.show()
```



[]:



hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

Conclusion:

1. Accuracy is a measure of overall correctness and is relatively high at approximately 87%, indicating that a significant portion of predictions is accurate.

The confusion matrix provides more detailed information about the model's performance. It shows that there are 6351 true negatives, 1483 true positives, 815 false positives, and 402 false negatives. This information helps in understanding how the model performs with respect to each class.

Precision is the ratio of true positives to the total predicted positives (true positives + false positives). The precision for class 1 (positive class) is approximately 0.79, indicating that when the model predicts a positive outcome, it is correct about 79% of the time.

Recall (or sensitivity) is the ratio of true positives to the total actual positives (true positives + false negatives). The recall for class 1 is approximately 0.65, which means the model correctly identifies about 65% of all actual positive instances.

F1-score is the harmonic mean of precision and recall and provides a balance between the two. The F1-score for class 1 is approximately 0.71, reflecting the trade-off between precision and recall.

2. The trade-offs must be taken into account when contrasting the outcomes of using the boosting and random forest algorithms on the Adult Census Income Dataset. Although there may be some interpretability trade-offs, boosting typically offers improved forecast accuracy, particularly for complex datasets. While maintaining



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

better interpretability and durability to overfitting, random forests, on the other hand, provide accuracy that is competitive.