

Experiment No. 3
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance: 16/08/2023
Date of Submission: 13/09/2023

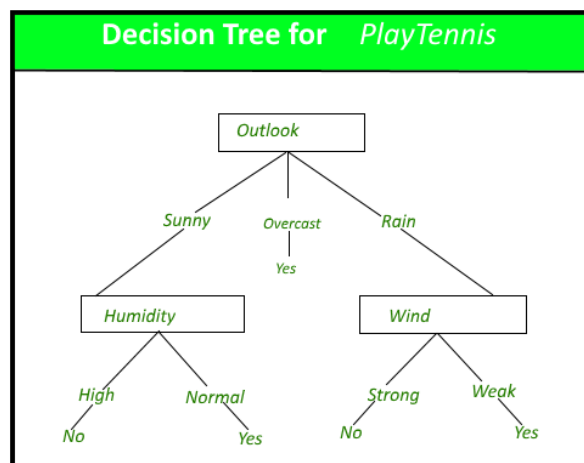


Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala,



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Code:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import numpy as np
```

```
df = pd.read_csv('./adult.csv')
```

```
df.describe()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

```
df.head(5)
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.loss
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	0	0
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	0	0
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	0	0
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	0	0
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	0	0

```
print ("Rows : ",df.shape[0])
print ("Columns : ",df.shape[1])
print ("\nFeatures : \n" ,df.columns.tolist())
print ("\nMissing values : ", df.isnull().sum().values.sum())
print ("\nUnique values : \n",df.nunique())
```

```
Rows : 32561
Columns : 15
```

```
Features :
['age', 'workclass', 'fnlwgt', 'education', 'education.num', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'capital.gain', 'capital.loss']
```

```
Missing values : 0
```

```
Unique values :
age          73
workclass     9
fnlwgt       21648
education     16
education.num 16
marital.status 7
occupation    15
relationship   6
race          5
sex           2
capital.gain  119
capital.loss   92
hours.per.week 94
native.country 42
income        2
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
```

```
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age          32561 non-null   int64
1   workclass    32561 non-null   object
2   fnlwgt       32561 non-null   int64
3   education    32561 non-null   object
4   education.num 32561 non-null   int64
5   marital.status 32561 non-null   object
6   occupation   32561 non-null   object
7   relationship  32561 non-null   object
8   race         32561 non-null   object
9   sex          32561 non-null   object
10  capital.gain  32561 non-null   int64
11  capital.loss  32561 non-null   int64
12  hours.per.week 32561 non-null   int64
13  native.country 32561 non-null   object
14  income       32561 non-null   object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
# checking "?" total values present in particular 'workclass' feature
df_check_missing_workclass = (df['workclass']=='?').sum()
df_check_missing_workclass
```

```
1836
```

```
# checking "?" total values present in particular 'occupation' feature
df_check_missing_occupation = (df['occupation']=='?').sum()
df_check_missing_occupation
```

```
1843
```

```
# checking "?" values, how many are there in the whole dataset
df_missing = (df=='?').sum()
df_missing
```

```
age          0
workclass    1836
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   1843
relationship  0
race         0
sex          0
capital.gain  0
capital.loss  0
hours.per.week 0
native.country 583
income       0
dtype: int64
```

```
percent_missing = (df=='?').sum() * 100/len(df)
percent_missing
```

```
age          0.000000
workclass    5.638647
fnlwgt       0.000000
education    0.000000
education.num 0.000000
marital.status 0.000000
occupation   5.660146
relationship  0.000000
race         0.000000
sex          0.000000
capital.gain  0.000000
capital.loss  0.000000
hours.per.week 0.000000
native.country 1.790486
income       0.000000
dtype: float64
```

```
# find total number of rows which doesn't contain any missing value as '?'
df.apply(lambda x: x != '?',axis=1).sum()
```

```
age          32561
workclass    30725
fnlwgt       32561
education    32561
education.num 32561
```

```

marital.status    32561
occupation        30718
relationship      32561
race              32561
sex               32561
capital.gain      32561
capital.loss      32561
hours.per.week    32561
native.country    31978
income            32561
dtype: int64

```

```
# dropping the rows having missing values in workclass
```

```
df = df[df['workclass'] != '?']
```

```
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	ca
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female		0
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female		0
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female		0
5	34	Private	216864	HS-grad	9	Divorced	Other-service	Unmarried	White	Female		0
6	38	Private	150601	10th	6	Separated	Adm-clerical	Unmarried	White	Male		0

```
# select all categorical variables
```

```
df_categorical = df.select_dtypes(include=['object'])
```

```
# checking whether any other column contains '?' value
```

```
df_categorical.apply(lambda x: x=='?',axis=1).sum()
```

```

workclass        0
education         0
marital.status    0
occupation        7
relationship      0
race              0
sex               0
native.country    556
income            0
dtype: int64

```

```
# dropping the "?"s from occupation and native.country
```

```
df = df[df['occupation'] != '?']
```

```
df = df[df['native.country'] != '?']
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age              30162 non-null  int64
1   workclass        30162 non-null  object
2   fnlwgt           30162 non-null  int64
3   education         30162 non-null  object
4   education.num     30162 non-null  int64
5   marital.status    30162 non-null  object
6   occupation        30162 non-null  object
7   relationship      30162 non-null  object
8   race              30162 non-null  object
9   sex              30162 non-null  object
10  capital.gain      30162 non-null  int64
11  capital.loss      30162 non-null  int64
12  hours.per.week    30162 non-null  int64
13  native.country    30162 non-null  object
14  income            30162 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```
from sklearn import preprocessing
```

```
# encode categorical variables using label Encoder
```

```
# select all categorical variables
```

```
df_categorical = df.select_dtypes(include=['object'])
```

```
df_categorical.head()
```

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
1	Private	HS-grad	Widowed	Exec-managerial	Not-in-family	White	Female	United-States	<=50K
3	Private	7th-8th	Divorced	Machine-op-inspct	Unmarried	White	Female	United-States	<=50K
4	Private	Some-college	Separated	Prof-specialty	Own-child	White	Female	United-States	<=50K
5	Private	HS-grad	Divorced	Other-service	Unmarried	White	Female	United-States	<=50K
6	Private	10th	Separated	Adm-clerical	Unmarried	White	Male	United-States	<=50K

```
# apply label encoder to df_categorical
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
1	2	11	6	3	1	4	0	38	0
3	2	5	0	6	4	4	0	38	0
4	2	15	5	9	3	4	0	38	0
5	2	11	0	7	4	4	0	38	0
6	2	0	5	0	4	4	1	38	0

```
# Next, Concatenate df_categorical dataframe with original df (dataframe)
# first, Drop earlier duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	marital.status	occupation	relation
1	82	132870	9	0	4356	18	2	11	6	3	
3	54	140359	4	0	3900	40	2	5	0	6	
4	41	264663	10	0	3900	40	2	15	5	9	
5	34	216864	9	0	3770	45	2	11	0	7	
6	38	150601	6	0	3770	40	2	0	5	0	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                 30162 non-null  int64
1   fnlwgt              30162 non-null  int64
2   education.num       30162 non-null  int64
3   capital.gain        30162 non-null  int64
4   capital.loss        30162 non-null  int64
5   hours.per.week      30162 non-null  int64
6   workclass           30162 non-null  int64
7   education           30162 non-null  int64
8   marital.status      30162 non-null  int64
9   occupation          30162 non-null  int64
10  relationship        30162 non-null  int64
11  race                30162 non-null  int64
12  sex                 30162 non-null  int64
13  native.country      30162 non-null  int64
14  income              30162 non-null  int64
dtypes: int64(15)
memory usage: 3.7 MB
```

```
# convert target variable income to categorical
df['income'] = df['income'].astype('category')
# check df info again whether everything is in right format or not
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                 30162 non-null  int64
```



```

1  fnlwgt      30162 non-null  int64
2  education.num  30162 non-null  int64
3  capital.gain  30162 non-null  int64
4  capital.loss  30162 non-null  int64
5  hours.per.week  30162 non-null  int64
6  workclass     30162 non-null  int64
7  education     30162 non-null  int64
8  marital.status 30162 non-null  int64
9  occupation    30162 non-null  int64
10 relationship  30162 non-null  int64
11 race         30162 non-null  int64
12 sex          30162 non-null  int64
13 native.country 30162 non-null  int64
14 income       30162 non-null  category
dtypes: category(1), int64(14)
memory usage: 3.5 MB

```

```

# Importing train_test_split
from sklearn.model_selection import train_test_split
# Putting independent variables/features to X
X = df.drop('income',axis=1)
# Putting response/dependent variable/feature to y
y = df['income']

```

```
X.head(5)
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	marital.status	occupation	relationship
1	82	132870	9	0	4356	18	2	11	6	3	
3	54	140359	4	0	3900	40	2	5	0	6	
4	41	264663	10	0	3900	40	2	15	5	9	
5	34	216864	9	0	3770	45	2	11	0	7	
6	38	150601	6	0	3770	40	2	0	5	0	

```
y.head(3)
```

```

1    0
3    0
4    0
Name: income, dtype: category
Categories (2, int64): [0, 1]

```

```

# Splitting the data into train and test
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.30,random_state=99)
X_train.head()

```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	marital.status	occupation	relationship
24351	42	289636	9	0	0	46	2	11	2	13	
15626	37	52465	9	0	0	40	1	11	4	7	
4347	38	125933	14	0	0	40	0	12	2	9	
23972	44	183829	13	0	0	38	5	9	4	0	
26843	35	198841	11	0	0	35	2	8	0	12	

```

# Importing decision tree classifier from sklearn library
from sklearn.tree import DecisionTreeClassifier
# Fitting the decision tree with default hyperparameters, apart from
# max_depth which is 5 so that we can plot and read the tree.
dt_default = DecisionTreeClassifier(max_depth=5)
dt_default.fit(X_train,y_train)

```

```

DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5)

```

```

# check the evaluation metrics of our default model
# Importing classification report and confusion matrix from sklearn metrics
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
# making predictions
y_pred_default = dt_default.predict(X_test)
# Printing classifier report after prediction
print(classification_report(y_test,y_pred_default))

```

	precision	recall	f1-score	support
0	0.86	0.95	0.91	6867
1	0.78	0.52	0.63	2182
accuracy			0.85	9049
macro avg	0.82	0.74	0.77	9049
weighted avg	0.84	0.85	0.84	9049

```
# Printing confusion matrix and accuracy
print(confusion_matrix(y_test,y_pred_default))
print(accuracy_score(y_test,y_pred_default))
```

```
[[6553 314]
 [1039 1143]]
0.8504807161012267
```

```
!pip install my-package
```

```
Collecting my-package
  Downloading my_package-0.0.0-py3-none-any.whl (2.0 kB)
Installing collected packages: my-package
Successfully installed my-package-0.0.0
```

```
!pip install pydotplus
```

```
Requirement already satisfied: pydotplus in /usr/local/lib/python3.10/dist-packages (2.0.2)
Requirement already satisfied: pyparsing>=2.0.1 in /usr/local/lib/python3.10/dist-packages (from pydotplus) (3.1.1)
```

```
# Importing required packages for visualization
from IPython.display import Image
from six import StringIO
from sklearn.tree import export_graphviz
import pydotplus,graphviz
# Putting features
features = list(df.columns[1:])
features
```


```
['fnlwgt',
 'education.num',
 'capital.gain',
 'capital.loss',
 'hours.per.week',
 'workclass',
 'education',
 'marital.status',
 'occupation',
 'relationship',
 'race',
 'sex',
 'native.country',
 'income']
```

```
!pip install graphviz
```

```
Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (0.20.1)
```

```
# plotting tree with max_depth=3
dot_data = StringIO()
export_graphviz(dt_default, out_file=dot_data,
feature_names=features, filled=True,rounded=True)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```

```
# GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
# specify number of folds for k-fold CV
n_folds = 5
# parameters to build the model on
parameters = {'max_depth': range(1, 40)}
# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
random_state = 100)
# fit tree on training data
tree = GridSearchCV(dtree, parameters,
cv=n_folds,
scoring="accuracy")
tree.fit(X_train, y_train)
```



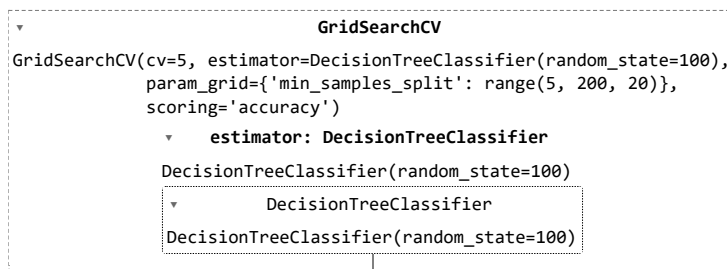
```
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	params	split0_test_score	split1_test_score
0	0.018425	0.001211	0.005323	0.000191	1	{'max_depth': 1}	0.747810	0.747810
1	0.035381	0.016654	0.007002	0.003257	2	{'max_depth': 2}	0.812219	0.818612
2	0.048566	0.017279	0.008126	0.003305	3	{'max_depth': 3}	0.828558	0.834241
3	0.040141	0.018838	0.003873	0.000468	4	{'max_depth': 4}	0.832583	0.840871
4	0.073534	0.036241	0.009387	0.008625	5	{'max_depth': 5}	0.834241	0.844897

```
# GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
# specify number of folds for k-fold CV
n_folds = 5
# parameters to build the model on
parameters = {'min_samples_leaf': range(5, 200, 20)}
# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
random_state = 100)
# fit tree on training data
tree = GridSearchCV(dtree, parameters,
cv=n_folds,
scoring="accuracy")
tree.fit(X_train, y_train)
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_min_samples_leaf	params	split0_test_score	split:
0	0.129880	0.012816	0.008649	0.003510	5	{'min_samples_leaf': 5}	0.825716	

```
# GridSearchCV to find optimal min_samples_split
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
# specify number of folds for k-fold CV
n_folds = 5
# parameters to build the model on
parameters = {'min_samples_split': range(5, 200, 20)}
# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
random_state = 100)
# fit tree on training data
tree = GridSearchCV(dtree, parameters,
cv=n_folds,
scoring="accuracy")
tree.fit(X_train, y_train)
```



```
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_min_samples_split	params	split0_test_score	spli
0	0.165695	0.022669	0.010115	0.004717	5	{'min_samples_split': 5}	0.811982	
1	0.194452	0.047568	0.008605	0.005247	25	{'min_samples_split': 25}	0.825006	
2	0.125775	0.006459	0.006552	0.000275	45	{'min_samples_split': 45}	0.835188	
3	0.128474	0.019117	0.006556	0.000719	65	{'min_samples_split': 65}	0.839451	
4	0.120904	0.010535	0.006089	0.000474	85	{'min_samples_split': 85}	0.846081	

```
# Create the parameter grid
param_grid = {
'max_depth': range(5, 15, 5),
'min_samples_leaf': range(50, 150, 50),
'min_samples_split': range(50, 150, 50),
'criterion': ["entropy", "gini"]
}
n_folds = 5
# Instantiate the grid search model
dtree = DecisionTreeClassifier()
grid_search = GridSearchCV(estimator = dtree, param_grid = param_grid,
cv = n_folds, verbose = 1)
# Fit the grid search to the data
grid_search.fit(X_train,y_train)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

GridSearchCV

```
GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),  
             param_grid={'criterion': ['entropy', 'gini']})
```

cv results

```
cv_results = pd.DataFrame(grid_search.cv_results_)
```

```
verbose=1)
```

cv_results

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_criterion	param_max_depth	param_min_samples_leaf	param_mi
0	0.109761	0.027735	0.007429	0.003987	entropy	5	50	
1	0.077615	0.021292	0.007246	0.001316	entropy	5	50	
2	0.059589	0.004072	0.005806	0.000180	entropy	5	100	
3	0.064347	0.011536	0.005823	0.000098	entropy	5	100	
4	0.091334	0.007191	0.005694	0.000245	entropy	10	50	
5	0.070263	0.013116	0.003854	0.000769	entropy	10	50	
6	0.058181	0.003006	0.004055	0.000751	entropy	10	100	



Conclusion:

1. Discuss about how categorical attributes have been dealt with during data pre-processing.



In data preprocessing, categorical attributes are typically converted into numerical representations to make them suitable for machine learning algorithms which is achieved through methods such as label encoding. Label encoding assigns a unique integer to each category. Additionally, we also choose to drop a column containing categorical values since it doesn't provide meaningful information for the analysis.

2. Discuss the hyper-parameter tuning done based on the decision tree obtained.



i) Max Depth: By limiting the decision tree's depth, this parameter keeps it from overcomplicating and overfitting the training set of data.

ii) Min Samples Split: This setting establishes the minimal amount of samples needed in a node to be subject to additional splitting. It aids in reducing the tree's tendency to specified choices depending on a constrained set of circumstances.

ii) Min Samples Leaf: This setting determines how many samples must be present in a leaf node. This can stop the tree from forming nodes with the same properties as min samples split extremely few occasions.

iv) Criteria: This option specifies the method for calculating a split's quality. Common criteria include "entropy" and "impurity".



3. Comment on the accuracy, confusion matrix, precision, recall and F1 score obtained.

→

- Accuracy is a measure of overall correctness and is relatively high at approximately 85.05%, indicating that a significant portion of predictions is accurate.
- The confusion matrix provides more detailed information about the model's performance. It shows that there are 6553 true negatives, 1143 true positives, 1039 false positives, and 314 false negatives. This information helps in understanding how the model performs with respect to each class.
- Precision is the ratio of true positives to the total predicted positives (true positives + false positives). The precision for class 1 (positive class) is approximately 0.78, indicating that when the model predicts a positive outcome, it is correct about 78% of the time.
- Recall (or sensitivity) is the ratio of true positives to the total actual positives (true positives + false negatives). The recall for class 1 is approximately 0.52, which means the model correctly identifies about 52% of all actual positive instances.
- F1-score is the harmonic mean of precision and recall and provides a balance between the two. The F1-score for class 1 is approximately 0.63, reflecting the trade-off between precision and recall.