



Experiment No. 4
Apply Random Forest Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:16-08-2023
Date of Submission:27-09-2023



Aim: Apply Random Forest Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Able to perform various feature engineering tasks, apply Random Forest Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

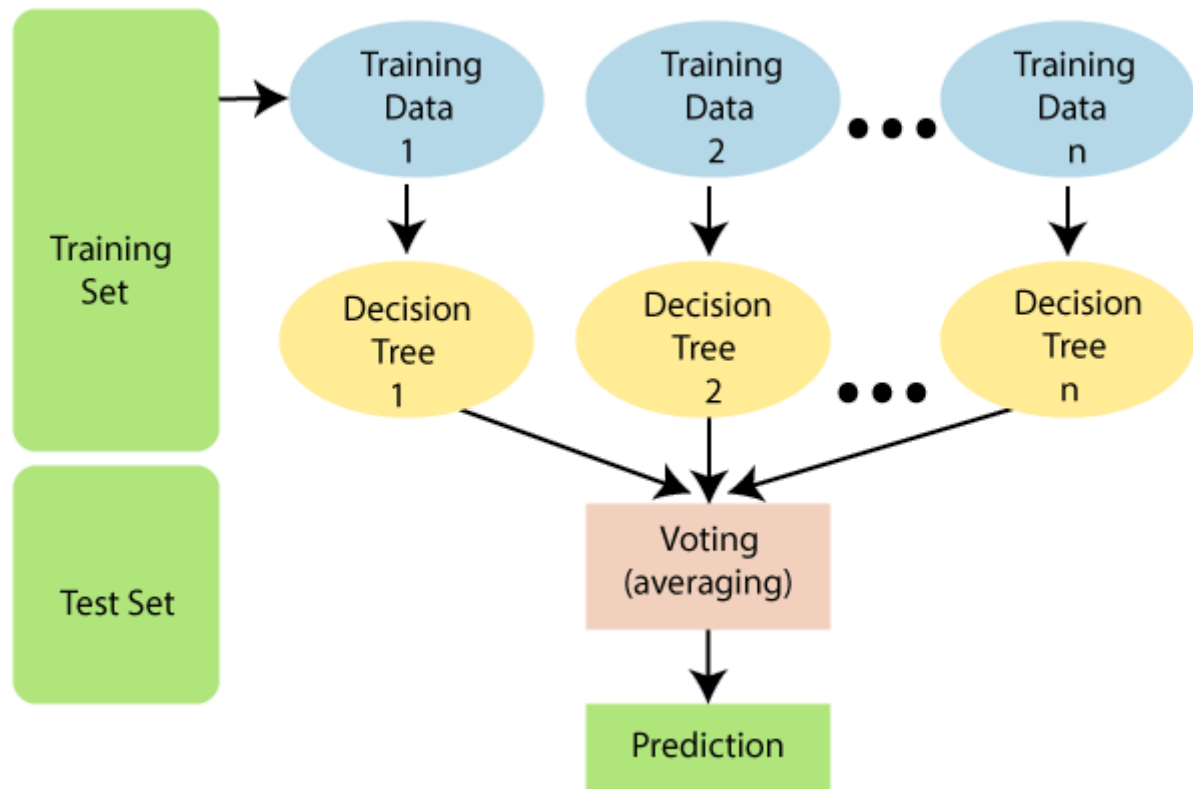
Theory:

Random Forest is a popular machine learning algorithm that belongs to the supervised learning technique. It can be used for both Classification and Regression problems in ML. It is based on the concept of ensemble learning, which is a process of combining multiple classifiers to solve a complex problem and to improve the performance of the model.

As the name suggests, "Random Forest is a classifier that contains a number of decision trees on various subsets of the given dataset and takes the average to improve the predictive accuracy of that dataset." Instead of relying on one decision tree, the random forest takes the prediction from each tree and based on the majority votes of predictions, and it predicts the final output.

The greater number of trees in the forest leads to higher accuracy and prevents the problem of overfitting.

The below diagram explains the working of the Random Forest algorithm:



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain:

continuous. capital-loss:

continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

exp-4-ml

October 9, 2023

```
[ ]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set(style='white', context='notebook', palette='deep')
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, cross_val_score, \
    StratifiedKFold, learning_curve, train_test_split, KFold
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
# To ignore warning messages
import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: df = pd.read_csv('./adult.csv')
```

```
[ ]: df.head(5)
```

```
[ ]:
  age workclass  fnlwgt   education  education.num marital.status \
0   90         ?   77053     HS-grad             9         Widowed
1   82   Private  132870     HS-grad             9         Widowed
2   66         ?  186061  Some-college            10         Widowed
3   54   Private  140359       7th-8th             4         Divorced
4   41   Private  264663  Some-college            10         Separated

   occupation  relationship   race   sex  capital.gain \
0           ?  Not-in-family  White  Female             0
1  Exec-managerial  Not-in-family  White  Female             0
2           ?    Unmarried  Black  Female             0
3  Machine-op-inspct    Unmarried  White  Female             0
4   Prof-specialty    Own-child  White  Female             0

  capital.loss  hours.per.week  native.country  income
```

0	4356	40	United-States	<=50K
1	4356	18	United-States	<=50K
2	4356	40	United-States	<=50K
3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K

```
[ ]: print ("Rows : " ,df.shape[0])
      print ("Columns : " ,df.shape[1])
      print ("\nFeatures : \n" ,df.columns.tolist())
      print ("\nMissing values : ", df.isnull().sum().values.sum())
      print ("\nUnique values : \n",df.nunique())
```

```
Rows : 32561
Columns : 15
```

```
Features :
['age', 'workclass', 'fnlwgt', 'education', 'education.num', 'marital.status',
'occupation', 'relationship', 'race', 'sex', 'capital.gain', 'capital.loss',
'hours.per.week', 'native.country', 'income']
```

```
Missing values : 0
```

```
Unique values :
age          73
workclass     9
fnlwgt      21648
education    16
education.num 16
marital.status 7
occupation   15
relationship  6
race         5
sex          2
capital.gain 119
capital.loss  92
hours.per.week 94
native.country 42
income       2
dtype: int64
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             32561 non-null  int64
```

```

1  workclass      32561 non-null  object
2  fnlwgt         32561 non-null  int64
3  education      32561 non-null  object
4  education.num  32561 non-null  int64
5  marital.status 32561 non-null  object
6  occupation     32561 non-null  object
7  relationship   32561 non-null  object
8  race           32561 non-null  object
9  sex            32561 non-null  object
10 capital.gain   32561 non-null  int64
11 capital.loss   32561 non-null  int64
12 hours.per.week 32561 non-null  int64
13 native.country 32561 non-null  object
14 income         32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```
[ ]: df.describe()
```

```

[ ]:
count      32561.000000  3.256100e+04  32561.000000  32561.000000  32561.000000  \
mean         38.581647  1.897784e+05    10.080679    1077.648844     87.303830
std          13.640433  1.055500e+05     2.572720    7385.292085    402.960219
min          17.000000  1.228500e+04     1.000000     0.000000     0.000000
25%          28.000000  1.178270e+05     9.000000     0.000000     0.000000
50%          37.000000  1.783560e+05    10.000000     0.000000     0.000000
75%          48.000000  2.370510e+05    12.000000     0.000000     0.000000
max          90.000000  1.484705e+06    16.000000   99999.000000   4356.000000

      hours.per.week
count      32561.000000
mean         40.437456
std          12.347429
min           1.000000
25%          40.000000
50%          40.000000
75%          45.000000
max          99.000000

```

```

[ ]: # checking "?" total values present in particular 'workclass' feature
df_check_missing_workclass = (df['workclass']=='?').sum()
df_check_missing_workclass

```

```
[ ]: 1836
```

```

[ ]: # checking "?" total values present in particular 'occupation' feature
df_check_missing_occupation = (df['occupation']=='?').sum()

```

```
df_check_missing_occupation
```

```
[ ]: 1843
```

```
[ ]: # checking "?" values, how many are there in the whole dataset  
df_missing = (df=='?').sum()  
df_missing
```

```
[ ]: age                0  
workclass             1836  
fnlwgt                0  
education             0  
education.num         0  
marital.status        0  
occupation            1843  
relationship          0  
race                  0  
sex                   0  
capital.gain          0  
capital.loss          0  
hours.per.week        0  
native.country        583  
income                0  
dtype: int64
```

```
[ ]: percent_missing = (df=='?').sum() * 100/len(df)  
percent_missing
```

```
[ ]: age                0.000000  
workclass             5.638647  
fnlwgt                0.000000  
education             0.000000  
education.num         0.000000  
marital.status        0.000000  
occupation            5.660146  
relationship          0.000000  
race                  0.000000  
sex                   0.000000  
capital.gain          0.000000  
capital.loss          0.000000  
hours.per.week        0.000000  
native.country        1.790486  
income                0.000000  
dtype: float64
```

```
[ ]: # find total number of rows which doesn't contain any missing value as '?'  
df.apply(lambda x: x != '?',axis=1).sum()
```



```
[ ]: age          32561
     workclass    30725
     fnlwgt       32561
     education    32561
     education.num 32561
     marital.status 32561
     occupation   30718
     relationship 32561
     race         32561
     sex          32561
     capital.gain 32561
     capital.loss 32561
     hours.per.week 32561
     native.country 31978
     income       32561
     dtype: int64
```

```
[ ]: # dropping the rows having missing values in workclass
df = df[df['workclass'] != '?']
df.head(5)
```

```
[ ]:   age workclass  fnlwgt   education  education.num marital.status \
1   82   Private  132870     HS-grad           9      Widowed
3   54   Private  140359     7th-8th           4      Divorced
4   41   Private  264663  Some-college          10      Separated
5   34   Private  216864     HS-grad           9      Divorced
6   38   Private  150601       10th           6      Separated
```

```
      occupation  relationship  race  sex  capital.gain \
1  Exec-managerial  Not-in-family  White  Female          0
3  Machine-op-inspct  Unmarried  White  Female          0
4   Prof-specialty  Own-child  White  Female          0
5   Other-service  Unmarried  White  Female          0
6   Adm-clerical  Unmarried  White  Male          0
```

```
      capital.loss  hours.per.week  native.country  income
1           4356           18  United-States  <=50K
3           3900           40  United-States  <=50K
4           3900           40  United-States  <=50K
5           3770           45  United-States  <=50K
6           3770           40  United-States  <=50K
```

```
[ ]: # select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
# checking whether any other column contains '?' value
df_categorical.apply(lambda x: x=='?',axis=1).sum()
```

```
[ ]: workclass      0
     education      0
     marital.status  0
     occupation      7
     relationship    0
     race            0
     sex             0
     native.country  556
     income          0
     dtype: int64
```

```
[ ]: # dropping the "?"s from occupation and native.country
df = df[df['occupation'] != '?']
df = df[df['native.country'] != '?']
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             30162 non-null  int64
1   workclass       30162 non-null  object
2   fnlwgt         30162 non-null  int64
3   education       30162 non-null  object
4   education.num   30162 non-null  int64
5   marital.status  30162 non-null  object
6   occupation      30162 non-null  object
7   relationship    30162 non-null  object
8   race            30162 non-null  object
9   sex             30162 non-null  object
10  capital.gain    30162 non-null  int64
11  capital.loss    30162 non-null  int64
12  hours.per.week  30162 non-null  int64
13  native.country  30162 non-null  object
14  income          30162 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
[ ]: from sklearn import preprocessing
     # encode categorical variables using label Encoder
     # select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()
```

```
[ ]: workclass      education marital.status      occupation  relationship \
1   Private        HS-grad      Widowed      Exec-managerial  Not-in-family
```

3	Private	7th-8th	Divorced	Machine-op-inspct	Unmarried
4	Private	Some-college	Separated	Prof-specialty	Own-child
5	Private	HS-grad	Divorced	Other-service	Unmarried
6	Private	10th	Separated	Adm-clerical	Unmarried

	race	sex	native.country	income
1	White	Female	United-States	<=50K
3	White	Female	United-States	<=50K
4	White	Female	United-States	<=50K
5	White	Female	United-States	<=50K
6	White	Male	United-States	<=50K

```
[ ]: # apply label encoder to df_categorical
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

```
[ ]: workclass  education  marital.status  occupation  relationship  race  sex \
1          2         11           6           3           1     4    0
3          2          5           0           6           4     4    0
4          2         15           5           9           3     4    0
5          2         11           0           7           4     4    0
6          2          0           5           0           4     4    1
```

	native.country	income
1	38	0
3	38	0
4	38	0
5	38	0
6	38	0

```
[ ]: # Next, Concatenate df_categorical dataframe with original df (dataframe)
# first, Drop earlier duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head(5)
```

```
[ ]: age  fnlwgt  education.num  capital.gain  capital.loss  hours.per.week \
1   82  132870           9           0         4356           18
3   54  140359           4           0         3900           40
4   41  264663          10           0         3900           40
5   34  216864           9           0         3770           45
6   38  150601           6           0         3770           40

workclass  education  marital.status  occupation  relationship  race  sex \
1          2         11           6           3           1     4    0
3          2          5           0           6           4     4    0
```

4	2	15	5	9	3	4	0
5	2	11	0	7	4	4	0
6	2	0	5	0	4	4	1

	native.country	income
1	38	0
3	38	0
4	38	0
5	38	0
6	38	0

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   30162 non-null  int64
1   fnlwt                 30162 non-null  int64
2   education.num         30162 non-null  int64
3   capital.gain           30162 non-null  int64
4   capital.loss           30162 non-null  int64
5   hours.per.week         30162 non-null  int64
6   workclass              30162 non-null  int64
7   education              30162 non-null  int64
8   marital.status         30162 non-null  int64
9   occupation             30162 non-null  int64
10  relationship           30162 non-null  int64
11  race                   30162 non-null  int64
12  sex                    30162 non-null  int64
13  native.country         30162 non-null  int64
14  income                 30162 non-null  int64
dtypes: int64(15)
memory usage: 3.7 MB
```

```
[ ]: plt.figure(figsize=(14,10))
sns.heatmap(df.corr(),annot=True,fmt='.2f')
plt.show()
```



```
[ ]: # convert target variable income to categorical
df['income'] = df['income'].astype('category')
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   30162 non-null  int64
1   fnlwgt                30162 non-null  int64
2   education.num         30162 non-null  int64
3   capital.gain          30162 non-null  int64
4   capital.loss          30162 non-null  int64
5   hours.per.week        30162 non-null  int64
6   workclass             30162 non-null  int64
7   education             30162 non-null  int64
8   marital.status        30162 non-null  int64
```

```

9  occupation      30162 non-null  int64
10 relationship    30162 non-null  int64
11 race            30162 non-null  int64
12 sex             30162 non-null  int64
13 native.country  30162 non-null  int64
14 income          30162 non-null  category
dtypes: category(1), int64(14)
memory usage: 3.5 MB

```

```
[ ]: # Importing train_test_split
from sklearn.model_selection import train_test_split
```

```
[ ]: # Putting independent variables/features to X
X = df.drop('income',axis=1)
# Putting response/dependent variable/feature to y
y = df['income']
```

```
[ ]: X.head(5)
```

```
[ ]:
   age  fnlwgt  education.num  capital.gain  capital.loss  hours.per.week  \
1   82  132870             9             0          4356             18
3   54  140359             4             0          3900             40
4   41  264663            10             0          3900             40
5   34  216864             9             0          3770             45
6   38  150601             6             0          3770             40

   workclass  education  marital.status  occupation  relationship  race  sex  \
1           2          11              6           3              1    4    0
3           2           5              0           6              4    4    0
4           2          15              5           9              3    4    0
5           2          11              0           7              4    4    0
6           2           0              5           0              4    4    1

   native.country
1              38
3              38
4              38
5              38
6              38

```

```
[ ]: y.head(5)
```

```
[ ]:
1    0
3    0
4    0
5    0
6    0

```

```
Name: income, dtype: category
Categories (2, int64): [0, 1]
```

```
[ ]: # Splitting the data into train and test
X_train,X_test,y_train,y_test = train_test_split(X,y)
X_train.head()
```

```
[ ]:      age  fnlwgt  education.num  capital.gain  capital.loss  hours.per.week  \
31795   48  207982             10             0             0             40
18861   63  113756             9             0             0             40
22081   37  405644             2             0             0             77
27727   43  240504            10             0             0             70
4200    43  397963             9           594             0             16

      workclass  education  marital.status  occupation  relationship  race  \
31795         2         15              4           7             4       2
18861         2         11              4           3             3       4
22081         2          3              3           4             2       4
27727         3         15              2           3             0       4
4200         2         11              0           5             1       4

      sex  native.country
31795   0              38
18861   0              38
22081   1              25
27727   1              38
4200    1              38
```

```
[ ]: test_size = 0.20
seed = 7
num_folds = 10
scoring = 'accuracy'
# Params for Random Forest
num_trees = 100
max_features = 3
models = []
```

```
[ ]: results = []
names = []
for name, model in models:
    kfold = KFold(n_splits=10, shuffle=True, random_state=seed)
    cv_results = cross_val_score(model, X_train, y_train, cv=kfold,
    ↪scoring='accuracy')
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
```

```
[ ]: random_forest = RandomForestClassifier(n_estimators=250,max_features=5)
random_forest.fit(X_train, y_train)
predictions = random_forest.predict(X_test)
print("Accuracy: %s%%" % (100*accuracy_score(y_test, predictions)))
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
```

Accuracy: 84.40525129293196%

[[5205 418]

[758 1160]]

	precision	recall	f1-score	support
0	0.87	0.93	0.90	5623
1	0.74	0.60	0.66	1918
accuracy			0.84	7541
macro avg	0.80	0.77	0.78	7541
weighted avg	0.84	0.84	0.84	7541

[]:



Conclusion:

A correlation heatmap is a graphical representation of a correlation matrix, where each cell in the heatmap represents the correlation between two variables. The correlation values are typically color-coded to help you quickly identify patterns. The correlation heatmap obtained from the dataset specifies significant positive correlations between education level and income, suggesting that higher education is associated with higher earnings.

Accuracy obtained in the decision tree model is 84.405%. A confusion matrix is a tabular representation used in machine learning to evaluate the performance of a classification model, especially for binary classification problems. Precision Obtain is 0.87 for 0 and 0.74 for 1, recall obtained is 0.93 for 0 and 0.60 for 1 and f1 score is 0.99 for 0 and 0.66 for 1.